



Part I

Final Report

FitML: A Fairness-Audited Machine Learning System for Sizing and Augmented Reality Try-On in E-Commerce



FitML: A Fairness-Audited Machine Learning System for Sizing and Augmented Reality Try-On in E-Commerce

Pillar 5: Capstone Project

Postgraduate Diploma in Artificial Intelligence and Machine Learning

Debie Galleros · July 2026

A virtual fitting room with size recommendations and tailored advice

Live demo: <https://fit-ml.netlify.app> · Repository:
<https://github.com/debiegalleros/fitml-capstone>

Contents

1 Problem Framing	3
2 Data Collection & Understanding	3
3 Preprocessing, EDA & Feature Engineering	4
4 Model Implementation & Comparison	7
5 Critical Thinking, Ethical AI & Bias Audit	8
6 The FitML Prototype: Try-On, Advice & Privacy	13
7 Presentations (Step 6)	14
8 GitHub Repository (Step 7)	14
9 Deployment (Step 8, optional)	15
10 Generative AI Usage (Step 9, optional)	15
11 Limitations & Future Work	16

1 Problem Framing

Business problem. Online clothing returns run at roughly 35%, versus about 8% in physical stores, costing retailers over \$100B annually. Sizing uncertainty is the leading driver, but shoppers also hesitate because they cannot see how a garment will actually look on their own body. FitML addresses both halves of that hesitation: a fairness-audited sizing model answers “will it fit?”, and a 2D augmented-reality try-on lets the shopper see the garment composited onto their own photo before buying — answering “how will this look on me?”. Color-suitability matching (does this shade complement the shopper) was identified as a natural extension and deliberately scoped out of this submission.

Data science problem. Supervised multi-class classification: given a customer's body measurements and the size they ordered, predict the reported fit outcome — **fit**, **runs small**, or **runs large** — paired with a fairness-audit layer that checks whether predictions are equally reliable across body types, heights, and size ranges. The labels come from real customer fit feedback, not collaborative filtering.

Success metrics. Technical: accuracy and macro-F1 (macro-F1 catches the minority-class failures that raw accuracy hides under a 73%-majority class), plus disparate impact ratio and equalized odds for the fairness layer. Business: projected return-rate reduction (35% toward a 15–20% target), with fewer wrong “it fits” reassurances as the mechanism.

The differentiator. Commercial sizing tools are black boxes. FitML publishes its full bias audit — per-group metrics, disparate impact, equalized odds, SHAP explainability, and an honest before/after mitigation comparison — as a first-class deliverable alongside the product.

2 Data Collection & Understanding

FitML deliberately combines sources with different, clearly-flagged statuses — real audited data where it exists, and a small synthetic extension where it does not:

Need	Source	Status
Model training + full fairness audit (graded core)	Kaggle “Clothing Fit Dataset for Size Recommendation” — 251,047 real customer reviews (58,503 ModCloth + 192,544 RentTheRunway) with self-reported measurements and fit feedback	Real
Men's catalog extension (~11 items, demo only)	200 synthetic customers labeled by rule-based lookup against Uniqlo's published men's size charts	Synthetic population, real sizing standard — excluded from the fairness audit
Garment catalog imagery (visual only, never training data)	Kaggle Fashion Product Images Dataset (MIT license) — 117 curated items, background-removed via rembg, 234 hue-shifted color variants	Real photos, demo pricing

Why this combination. Both ModCloth and RentTheRunway sell women's apparel, so the real dataset is women's-only. Rather than force one dataset to cover both genders — or silently train on invented men's fit feedback — the project uses real, citable data where it is available and a small, explicitly-documented synthetic extension for the gap. The men's extension is a deterministic size-chart lookup, not a trained model, and its ~200-row synthetic population is far too small for valid group comparisons — so it is excluded from the fairness audit by design, an explicit documented limitation rather than a silently patched one.

Key fields (full data dictionary in docs/data_dictionary.md). Height, weight, bust band + cup, hip, self-reported body type (hourglass / pear / apple / athletic / petite / straight & narrow / full bust), garment

category, ordered size, and the target **fit_feedback** (small / fit / large). Waist was dropped: RentTheRunway never collected it and ModCloth's own field was 96.5% null, leaving ~1% combined coverage — too sparse to impute defensibly. Post-purchase fields (rating, review text) were excluded as features because they would not exist at prediction time.

Catalog imagery provenance. iMaterialist Fashion was evaluated first and rejected: its images are in-the-wild street-style, red-carpet, and runway photos — busy backgrounds, identifiable people, some with third-party photographer watermarks — a poor fit both technically (for background-removal + affine-warp compositing) and ethically. The Fashion Product Images Dataset provides clean, plain-background product shots with structured gender/category/color metadata; only the ~115 needed full-resolution images were fetched via targeted single-file downloads, not the full 24.7 GB archive. Catalog prices are programmatically generated demo pricing within per-category PHP bands, documented as such.

3 Preprocessing, EDA & Feature Engineering

Full walkthrough with reasoning in notebooks/02_eda_feature_engineering.ipynb; the reusable logic lives in `src/prepare_data.py` so later phases import it instead of duplicating it. Output: `data/processed/model_ready.csv` — 251,047 rows × 15 columns, no remaining nulls.

2 Harmonizing two schemas

The sources disagree on units and formats (height as “5ft 6in” vs “5' 8””, hips in inches, bust as band+cup strings). Parsing normalizes everything to metric. ModCloth rows tagged only new/sale/wedding (24,287 rows, ~29% of ModCloth) carry no garment-category signal and were dropped; RentTheRunway's 68 fine-grained category tags were mapped onto ModCloth's 4-bucket taxonomy (tops / dresses / bottoms / outerwear) as `category_broad`. Cup letters were unified to an ordinal scale (AA=0 ... K=11) using the standard DD=E, DDD=F, DDDD=G equivalence.

Structural missingness, handled honestly

Missingness here is structural, not random: weight is absent for *all* ModCloth rows (never collected), hips for all RentTheRunway rows, body type for all ModCloth rows. Instead of silently imputing and pretending the values are measured, each such field gets a `_missing` indicator flag alongside median imputation, so models — and the fairness audit — can distinguish a measured value from an invented placeholder. `body_type` gets an explicit `not_reported` category rather than a guessed shape: inventing a protected-attribute-like label would contaminate the Phase 6 audit. This decision pays off later — SHAP shows the model actually uses these missingness flags (Section 5.3).

Class imbalance in the target




Figure 1. Target distribution: 72.7% fit, 13.7% large, 13.6% small.

The target is ~73/14/13 — exactly why macro-F1, not accuracy, drives model selection: a model that always predicts “fit” scores ~72.7% accuracy while being useless for the minority classes that actually need a size correction.

Distributions and correlations




Figure 2. Feature distributions per fit_feedback class. size_ordered separates most clearly; body measurements separate weakly on their own.




Figure 3. Correlation heatmap of numeric features.

size_ordered shows the clearest class separation (it is mechanically tied to fit), while body measurements separate weakly in isolation — fit depends on the relationship between measurements and the ordered size, not on measurements alone.

PCA and exploratory feature importance




Figure 4. PCA (2D, visualization only) colored by body type.

Body-type clusters overlap substantially in 2D: PC1/PC2 are dominated by overall body size, while body-shape categories are defined by proportional relationships (e.g. waist-to-hip ratio). This is a useful negative result — raw measurements will not trivially separate body types, reinforcing why the fairness audit needs per-group metrics rather than assuming a simple boundary.




Figure 5. Exploratory random-forest feature importances (fixed seed).

Leakage discipline. `model_ready.csv` ships deliberately unscaled and not one-hot encoded: scaling and encoding are fitted inside each Phase 5 model's train-only sklearn Pipeline, so no test-split statistics ever leak into training.

4 Model Implementation & Comparison

Four classifiers — Logistic Regression (baseline), Random Forest, XGBoost, and a two-hidden-layer MLP (64/32, early stopping) — were trained by `src/train_models.py` on the same stratified 80/20 split (200,837 train / 50,210 test rows, seed 42), same features, same untuned equal-footing hyperparameters, with preprocessing fitted inside each pipeline on the training fold only.

model	accuracy	macro-F1	F1 fit	F1 large	F1 small	train time (s)
xgboost	0.7272	0.2918	0.8418	0.0219	0.0115	2.7
random_forest	0.7278	0.2848	0.8424	0.0101	0.0018	3.4
mlp	0.7272	0.2825	0.8420	0.0043	0.0012	6.1
logistic_regression	0.7267	0.2822	0.8418	0.0000	0.0046	0.8

Table 1. Phase 5 comparison (`models/comparison.csv`). Selected: XGBoost, by macro-F1.

Why macro-F1

Accuracy is statistically indistinguishable across all four models — each lands at the 72.7% majority baseline — which is precisely why it was ruled out as the selection metric. Macro-F1 weights each class equally, exposing minority-class failure: the “small” and “large” predictions are the entire point of a sizing assistant, because they are the cases where a recommendation prevents a return. XGBoost wins as the only model with non-trivial F1 on both minority classes.

The honest headline: unweighted models collapse to the majority class

Every unweighted model learns to say “fit” almost always (logistic regression never predicts “large” at all). Two causes: (1) with 73% of labels being “fit”, an unweighted loss is minimized by rarely risking a minority prediction; (2) the profile features carry real but limited signal — fit feedback also depends on garment-specific cut and sizing quirks no body measurement captures, and published work on this same dataset reports the task as hard for the same reason. This was deliberate for Phase 5: the unweighted models are the honest “before” for the fairness mitigation in Section 5.

Trees vs. the MLP — the expected finding, reported honestly

The Master Plan predicted tree-based models would beat the MLP on this modest-size, mixed-type tabular dataset, and they did: the MLP lands below both XGBoost and Random Forest on macro-F1 despite the longest training time. This matches the well-documented pattern that gradient-boosted trees dominate MLPs on tabular data of this scale — a legitimate comparison result, not a tuning failure. CNNs and diffusion models were explicitly scoped out.

Which model the product serves

The deployed /recommend-size endpoint serves the **class-weighted XGBoost** from the Section 5 mitigation, not the unweighted Phase 5 winner. For a sizing assistant the costly error is a missed misfit — the unweighted model gives wrong “it fits” reassurance in ~96–100% of misfit cases, while the weighted model catches roughly half of misfits (recall large 0.486, small 0.546). The weighted model is also better on the selection metric itself (macro-F1 0.3596 vs 0.2918). Its accuracy cost (0.727 → 0.396) is absorbed by the product layer: the endpoint returns class probabilities, and the displayed confidence % plus a blue/amber advisory box communicate uncertainty, so borderline predictions surface as a sizing tip, never a bare verdict.

A deployment lesson: correlation is not causation

The first /recommend-size design scanned candidate sizes and picked the one maximizing P(fit) — and testing showed it inverts: a petite profile got XL at 90% “confidence”. Cause: in the training data, size_ordered tracks the customer's own body, so the model learned the between-customer correlation (larger sizes are ordered by larger people, who disproportionately report items running small), not the within-person effect of changing size. A (petite body, size 17) row is far out of distribution, and the model's counterfactual answer there is meaningless. The fix: anchor the size deterministically from measurements via the size chart, ask the model the in-distribution question (“does this size run small/fit/large for this body?” — literally what each training row records), shift one size if the top class is a misfit, then apply the borderline advisory layer. Verified sane: petite → XS, average → M, large-bodied → XL.

5 Critical Thinking, Ethical AI & Bias Audit

Full report in docs/fairness_report.md; every number below is generated by src/fairness_audit.py, which re-derives the exact Phase 5 split (seed 42) and asserts it matches the pinned test indices before computing anything. The audit covers the real women's model only — the men's synthetic extension is

excluded because group metrics on a ~200-row generated population would measure properties of the random-number generator, not real-world disparity.

5.1 Setup: harms, groups, metrics

A “fit” prediction is the favorable outcome (a confident go-ahead at the usual size); small/large predictions route the shopper into size-adjustment friction. Two harms matter: quality-of-service harm (misfit detection failing more for some groups → more wrong reassurance → more returns) and allocation-style harm (the favorable prediction handed out at unequal rates). Groups audited, each with $\geq 1,000$ test rows: **body_type** (8 groups incl. not_reported, $n=14,800$), **height_band** (under 160 / 160–169 / 170+ cm), and **size_band** (US 0–6 / 8–14 / 16–22 / 24+). Metrics: per-group accuracy and per-class recall, disparate impact ratio (four-fifths rule, <0.8 flagged), and equalized-odds TPR/FPR spreads.

5.2 Baseline audit: “fair” numbers that mean nothing

No group is DI-flagged at baseline (all ratios ≥ 0.98) — and that apparent fairness is vacuous. The unweighted model predicts “fit” for 99.5% of all rows in every group; a model that always says yes passes the four-fifths rule by construction. The metrics that look inside the predictions show real disparities:

- **Accuracy tracks body size:** size 0–6 shoppers get 0.7506, size 24+ get 0.6781 — a 7.3-point gap against plus-size shoppers, who experience misfit more often and get the least reliable model exactly there.
- **Misfit detection is near zero for everyone, but unevenly so:** full-bust shoppers get their “runs large” cases caught 8.2% of the time, pear-shaped shoppers 0.34% — a 24× ratio. The “small” class has literally zero recall for apple body types and size 0–6.
- **The fit-class false-positive rate is 0.96–1.00 in every group:** when an item does not fit, the model tells virtually every shopper that it does.

Conclusion: the baseline’s bias problem is not unequal favorable-outcome rates; it is a degenerate predictor whose failures concentrate on the shoppers most likely to experience misfit. The mitigation must attack minority-class recall.

5.3 Explainability: SHAP on the winning model




Figure 6. Global SHAP feature importance (TreeExplainer, 3,000 test rows, seed 42).

size_ordered dominates, especially for the minority classes (mean |SHAP| 0.38 for small, 0.27 for large vs 0.08 for fit), and the beeswarms read sensibly: ordering a high size, low weight, and a small bust band all push toward predicting “runs small”. The model has learned real, directionally correct sizing signal — it is the decision threshold, drowned by class imbalance, that fails.




Figure 7. SHAP beeswarm for the “small” class.

Fairness-relevant red flag: missingness is predictive. `weight_kg_missing` and `body_type_not_reported` rank in the top seven features for minority classes. These overlap heavily with `source_modcloth` (weight is missing for all ModCloth rows): the model partly keys on which platform the shopper came from and how much they disclosed, not on their body. Shoppers who decline to report body type or weight get systematically different predictions — a data-completeness bias the product answers with transparency (the confidence box lowers displayed confidence when key measurements are absent).

5.4 Mitigation: class-weighted retraining, before/after

One mitigation, fully audited, as scoped: identical pipeline, identical train fold, identical seed — the only change is balanced sample weights ($\text{weight} \propto 1/\text{class frequency}$) passed to XGBoost.




Figure 8. Per-class recall, before (unweighted) vs after (class-weighted).

metric	before (unweighted)	after (class-weighted)
macro-F1 (selection metric)	0.2918	0.3596
recall — large	0.0112	0.4864
recall — small	0.0059	0.5463
min per-group large recall (body_type)	0.0033	0.3553
min per-group small recall (body_type)	0.0000	0.4849
worst-group fit FPR (wrong “it fits”)	0.9957	0.3463
accuracy	0.7272	0.3959
DI flags (body_type + size_band)	0 (vacuous)	7 (informative)

Table 2. Mitigation before/after — the fairness-critical numbers.

Reading the trade honestly: misfit detection now exists for every group (the wrong-reassurance rate dropped from ~0.96–1.00 to 0.13–0.35 everywhere), and macro-F1 rose 23%. The costs are equally explicit: accuracy fell to 0.396, most “it fits” cases now receive a spurious caution, and — importantly — **disparate-impact flags appear**. With the model actually discriminating between classes, the favorable prediction goes to size 8–14 shoppers at rate 0.42 but size 24+ at 0.17 (DI 0.41, flagged), and apple / full-bust / not_reported body types flag at DI 0.64–0.70. Height stays equal (DI ≥ 0.99). These flags are informative, not an artifact: the baseline's tiny spreads were a side effect of predicting one class everywhere, not evidence of equal treatment. The weighted model partly reflects real, higher misfit base rates in those groups, but at ratios well past the 0.8 threshold and with spurious “runs small” warnings (FPR up to 0.60) for size 16+.

Deployment stance. The product serves the class-weighted model: catching misfits matters more than raw accuracy for a sizing assistant, and the confidence layer absorbs the uncertainty. A calibrated middle ground — probability calibration plus per-group decision thresholds tuned to equalize misfit-detection rates — is the clear next step, documented as future work to keep the before/after comparison clean.

5.5 Leakage checks and limitations

- Preprocessing fitted train-only inside pipelines; the audit asserts the re-derived split matches the pinned test indices; no post-purchase fields among features; size_ordered is legitimately known at prediction time.
- Residual risk, documented: the random split does not group by customer, so a repeat customer could appear on both sides. Customer IDs were dropped upstream; noted as a known limitation.
- Self-reported measurements are noisy, plausibly in group-correlated ways (social-desirability bias in weight reporting is well documented).
- Groups are proxies, not protected attributes — the data has no age, ethnicity, or disability fields; results must not be read as a full protected-attribute audit.
- The largest body_type group is “declined to say” (n=14,800), so per-group numbers for reported categories are conditioned on willingness to self-categorize.
- Bust input-method (band+cup native vs converted chest measurement) is a fairness-relevant data-quality axis the training data cannot audit; the product logs a per-user bust_input_method flag specifically so a future audit can stratify by it.

6 The FitML Prototype: Try-On, Advice & Privacy

A functional prototype website (not an online store — no checkout, payments, or inventory), live at <https://fit-ml.netlify.app> with six pages: Landing, About/How-it-works, Profile setup, Catalog, Try-on result, and History. Flask backend, plain HTML/CSS/JavaScript frontend (no framework), responsive from 375px mobile to wide desktop.

2D body-wrap try-on

MediaPipe extracts pose keypoints from the user's uploaded photo; a background-removed garment PNG is affine-warped (scaled and rotated to shoulder width and torso angle) and alpha-composited onto the photo on an HTML5 Canvas — genuine static-image AR in which the user's real photo, pose, and background remain fully visible. Full 3D (cloth simulation or diffusion-based rendering) was considered and ruled out as out of scope for a solo build; none of the AR fidelity is part of the graded rubric. When the shopper composites a size other than the recommended one, the garment is additionally scaled by that size's measurement ratios from the size chart — an XL on a small-framed user renders visibly wider and longer than an M — a lightweight visual size-difference cue, with the advice text explaining the fit tradeoff.

Size recommendation UX

Measurements are always manually entered (height, optional weight, bust via band+cup or converted chest measurement flagged lower-precision, hip, body type). The uploaded photo is for try-on pose only and never feeds the size model. The recommendation arrives in a confidence box: blue with a checkmark for confident matches; amber with a lightbulb “Sizing tip:” for borderline cases (near a size boundary + low-stretch fabric + fitted cut → recommend one size up, lower the displayed confidence). Advice text is two paragraphs: measurement-based reasoning, then a “Note:” paragraph of plain-language visual observations from the composite.

History-based suggestions (rule-based, not a second model)

Each try-on stores size shown, brand, fabric, and the user's fit feedback. When 2+ prior items from the same brand or fabric show consistent feedback, a one-line note is added to the advice prompt (“items from this brand have run true to size for you”) — a SQL lookup plus a conditional prompt line, deliberately

not another trained model.

Privacy (full framing in docs/privacy.md)

- Uploads live under anonymous UUID session folders (backend/uploads/{session_id}/), never name/email-keyed, gitignored, session-scoped.
- 24-hour auto-deletion enforced in code (purge on startup, before every upload, and as a scheduled script); the deployed host's ephemeral disk makes redeploys a stricter version of the same guarantee.
- Face blur ON by default: MediaPipe face detection + Gaussian blur applied before the photo is ever written to disk, so an unblurred version never persists unless the user explicitly opts in.
- HTTPS-only transmission; consent notice at the point of collection stating what is collected, why, and the retention period.
- Framed against the Philippines' Data Privacy Act of 2012 (RA 10173): transparency, legitimate purpose, proportionality, retention limitation, and security — with consent for the try-on feature never reused for research or auditing.

7 Presentations (Step 6)

Two decks accompany this report, both following the program branding (white background, AIM logo on every slide):

- **Technical deck** (Debie_Galleros_Pillar5_CapstoneProject_TechnicalDeck.pptx, audience: peers and faculty): problem framing and task type, dataset overview, EDA highlights, the four-model comparison with macro-F1 reasoning, fairness-audit methodology with before/after mitigation numbers, SHAP explainability, the data → model → Flask → website architecture, GenAI usage boundaries, and limitations + future work.
- **Business deck** (Debie_Galleros_Pillar5_CapstoneProject_BusinessDeck.pptx, audience: executives, no jargon): the returns problem (\$100B+, 35% online return rate), what FitML does in shopper language, the fairness differentiator, live-site screenshots, business impact and ROI logic, privacy and risk handling, rollout strategy, and future work.

8 GitHub Repository (Step 7)

<https://github.com/debiegalleros/fitml-capstone> — public, with the rubric-required structure (src/, notebooks/, data/, models/, backend/, frontend/, docs/, README.md, requirements.txt) and a clean, incremental commit history: one commit per meaningful unit of work with descriptive messages, phase by phase from scaffolding through deployment.

- Reproducibility: fixed seed 42 throughout; training config, label encoder, and pinned test indices committed alongside the model artifacts; the fairness audit refuses to run if the re-derived split stops matching the pinned indices.
- Hygiene: raw datasets, catalog image assets, uploaded photos, and API keys are gitignored; requirements.txt is frozen to the exact working versions; the one oversized artifact (random_forest.joblib, 50 MB) is gitignored and byte-reproducible from the fixed-seed training script.
- Documentation: eight docs/ files cover the data dictionary, model selection, the full fairness report, data provenance, privacy, GenAI usage, and deployment — all cross-linked from the README.

9 Deployment (Step 8, optional)

Live end-to-end on free tiers: the Flask backend runs as a Docker web service on Render (<https://fit-ml.onrender.com>) and the static frontend on Netlify (<https://fit-ml.netlify.app>). Free tiers spin down when idle — the first request after idle can take 30–60 seconds to cold-start.

- Docker was required because MediaPipe's Tasks API loads GPU-stack shared libraries (libGLv2) even on the CPU delegate; the Dockerfile installs the needed system libraries that Render's native Python runtime cannot.
- Catalog images are gitignored, so a build-time script fetches them from a GitHub Release asset — keeping binary blobs out of git history while making assets available on the ephemeral-disk host.
- The full flow was verified against the live backend: photo upload (with face blur and pose extraction) → size recommendation → try-on composite → Claude advice call.
- The ANTHROPIC_API_KEY lives only in the host's environment variables, never in the repository.

10 Generative AI Usage (Step 9, optional)

FitML uses generative / deep-learning AI in exactly two places, both deliberately outside the graded ML pipeline (full write-up in `docs/genai_usage.md`):

Claude API — multimodal fit advice

After the try-on composite renders, the `/advice` endpoint sends Claude the shopper's measurements, the model's recommendation and confidence, and the composited try-on image itself. Claude returns two paragraphs of plain-language advice that can reference what is visible in the composite. Hard boundaries: Claude never estimates body measurements (they are always manually entered), never overrides or feeds the trained size model (it receives the recommendation as read-only context), and its visual observations are qualitative commentary only. The prompt bans tailoring jargon so the output stays readable for non-technical shoppers.

rembg / U²-Net — catalog background removal

The catalog pipeline runs product photos through rembg's clothing-segmentation model (U²-Net, a pretrained CNN) to produce the transparent garment PNGs the compositor warps. It runs offline, once per image, sees no user data, and touches no training data.

Why GenAI is fenced off from the graded pipeline

- **Auditability:** the fairness audit makes quantitative claims about a model whose inputs and parameters are fixed and inspectable; a generative model in the prediction path would make those claims unverifiable.
- **No silent data invention:** letting a vision model estimate measurements from photos would inject unquantifiable, likely body-type-correlated error into exactly the variables the audit stratifies on.
- **Reproducibility:** every graded artifact reproduces from fixed seeds and committed code; generative output is confined to surfaces where variation is harmless (advice prose, one-time asset prep).

The result: deterministic, audited ML decides the size; generative AI only explains it and helps build the demo catalog.

11 Limitations & Future Work

- **Class imbalance remains the dominant modeling problem.** Next steps beyond the single audited mitigation: probability calibration with per-group decision thresholds, focal loss, or resampling.
- **Customer-grouped splitting:** re-derive customer IDs upstream so repeat customers cannot straddle the train/test split.
- **Photorealistic, fit-accurate try-on:** the FIT dataset (2026; 1.13M fit-aware try-on triplets) is the cited research path for generative try-on that renders how a specific size actually fits — the successor to this submission's size-proportional 2D scaling.
- **Transformer-based clothing parsing:** a SegFormer-B2-clothes benchmark against the current U²-Net pipeline found no net quality advantage today (59 vs 56 clean items, 5 outright failures), but the approach is the natural upgrade as models mature.
- **Color-suitability matching** (does this shade complement the shopper) — identified from the start as a natural extension, deliberately out of scope.
- **A CNN fit-quality assessor** over the composited try-on image was considered and scoped out: no labeled dataset exists, and building one from the compositing pipeline was too time/risk-heavy for the deadline.
- **Production hardening:** pin CORS to the exact frontend origin, add per-group threshold monitoring, and collect the bust_input_method flag long enough to audit input-method bias with real usage data.

References & data sources

- Kaggle: Clothing Fit Dataset for Size Recommendation (ModCloth + RentTheRunway) — model training and fairness audit.
- Kaggle: Fashion Product Images Dataset (paramaggarwal, MIT license) — catalog imagery.
- Uniqlo published men's size charts (transcribed 2026-07-06) — men's extension label source.
- Republic Act No. 10173 (Data Privacy Act of 2012, Philippines) — privacy framing.
- FIT dataset (2026): 1.13M fit-aware try-on triplets — future-work research direction.
- Project repository: <https://github.com/debiegalleros/fitml-capstone> · Live demo: <https://fit-ml.netlify.app>



Part II

Technical Presentation Deck

Audience: peers and faculty — models, fairness audit, architecture



FitML: A Fairness-Audited Machine Learning System for Sizing and Augmented Reality Try-On in E-Commerce

Pillar 5: Capstone Project

Postgraduate Diploma in Artificial Intelligence and Machine Learning

Debie Galleros · July 2026

A virtual fitting room with size recommendations and tailored advice

Technical presentation — fit-ml.netlify.app · github.com/debiegalleros/fitml-capstone



Problem framing and task type

35%

online clothing return rate (vs ~8% in physical stores)

\$100B+

annual cost of returns to retailers

2 questions

“will it fit?” and “how will it look on me?” — both drive hesitation

The data science task

- **Supervised multi-class classification:** predict reported fit outcome — fit / runs small / runs large — from body measurements + ordered size
- **Labels are real customer fit feedback** (not collaborative filtering)
- **Selection metric: macro-F1** — accuracy is blind to minority classes under a 73%-majority target

What makes it different

- **A published fairness audit** — disparate impact, equalized odds, SHAP, and a mitigation with honest before/after numbers
- Commercial sizing tools are black boxes; **FitML treats the audit as a first-class deliverable**
- **2D AR try-on** gives the visual answer without 3D or diffusion models (explicitly scoped out)



Dataset: 251,047 real fit reviews

Kaggle “Clothing Fit Dataset for Size Recommendation” — ModCloth + RentTheRunway


	Women's core (graded ML)	Men's extension (demo)
Customers	251,047 real reviews	200 synthetic
Labels	Real fit feedback	Uniqlo chart lookup
Sizing logic	4 trained classifiers	Deterministic rule
Fairness audit	Included (full)	Excluded (documented)

Why: both platforms sell women's apparel, so real data is women's-only. Real, citable data where it exists; a small, clearly-flagged synthetic extension for the gap — an explicit limitation, not a silently patched one.

Features (all self-reported / user-entered)

height · weight · bust band + cup · hip · body type · garment category · ordered size — target: fit / small / large

Class balance (the core difficulty)



72.7 / 13.7 / 13.6 — a model that always says “fit” gets 72.7% accuracy while being useless.



EDA and feature engineering

notebooks/02_eda_feature_engineering.ipynb · reusable logic in src/prepare_data.py

Two schemas, harmonized

Height parsed from “5ft 6in” vs “5' 8””; 68 RentTheRunway category tags mapped to a 4-bucket taxonomy; cup letters unified to an ordinal scale (AA=0 ... K=11). 24,287 ModCloth rows with no garment signal dropped.

Missingness is structural — flagged, not hidden

Weight absent for ALL ModCloth rows; hips for ALL RentTheRunway rows. Each gets a `_missing` indicator beside median imputation; `body_type` gets an explicit `not_reported` category (never a guessed shape). SHAP later proves the model uses these flags.

PCA: a useful negative result

Body-type clusters overlap heavily in 2D — components track overall size, not shape. Raw measurements won't trivially separate body types, so the fairness audit needs per-group metrics.

Leakage discipline

`model_ready.csv` ships unscaled and un-encoded; scaling and one-hot encoding are fitted inside each model's train-only Pipeline. Post-purchase fields (rating, review text) excluded as features.

Model comparison: four classifiers, one honest headline

Same stratified 80/20 split (200,837 / 50,210 rows, seed 42) · same features · preprocessing fitted train-only

model	accuracy	macro-F1	F1 fit	F1 large	F1 small
XGBoost (selected)	0.7272	0.2918	0.8418	0.0219	0.0115
Random Forest	0.7278	0.2848	0.8424	0.0101	0.0018
MLP (64/32)	0.7272	0.2825	0.8420	0.0043	0.0012
Logistic Regression	0.7267	0.2822	0.8418	0.0000	0.0046

Every unweighted model collapses toward the majority class — accuracy is pinned at the 72.7% baseline for all four; logistic regression never predicts “large” at all. XGBoost wins as the only model with non-trivial F1 on both minority classes.

Trees beat the MLP, as predicted — gradient-boosted trees dominate MLPs on modest tabular data. A legitimate finding, reported honestly (equal-footing, untuned hyperparameters for all four).

Why macro-F1?

“Small” and “large” are the whole point: they are the cases where a recommendation prevents a return.

Macro-F1 weights each class equally, exposing the minority-class failure that accuracy hides.

Accuracy would call a coin-flip between these models; macro-F1 separates them.



Fairness audit: methodology

src/fairness_audit.py re-derives the exact Phase 5 split and asserts it matches the pinned test indices

Two harms audited

- **Quality-of-service:** misfit detection failing more for some groups → wrong “it fits” reassurance → returns
- **Allocation-style:** the favorable “fit” prediction handed out at unequal rates

Groups (each $\geq 1,000$ test rows)

- **body_type** — 8 groups incl. not_reported (n=14,800)
- **height_band** — under 160 / 160–169 / 170+ cm
- **size_band** — US 0–6 / 8–14 / 16–22 / 24+

Metrics

Per-group accuracy and per-class recall · disparate impact ratio (four-fifths rule, <0.8 flagged) · equalized-odds TPR/FPR spreads

Baseline finding: “fair” numbers that mean nothing


No group is DI-flagged (all ≥ 0.98) — because the model predicts “fit” for 99.5% of rows in every group. A model that always says yes passes the four-fifths rule by construction.

Inside the predictions: a 7.3-point accuracy gap against size 24+ shoppers; “runs large” caught 8.2% of the time for full-bust vs 0.34% for pear (24×); zero “small” recall for apple body types.

The fit-class FPR is 0.96–1.00 everywhere: when an item does not fit, virtually every shopper is told it does.

Mitigation: class-weighted retraining — before / after

Identical pipeline, split, and seed; the only change is balanced sample weights




metric	before	after
macro-F1 (selection metric)	0.2918	0.3596
recall — large	0.011	0.486
recall — small	0.006	0.546
worst-group fit FPR	0.996	0.346
accuracy	0.727	0.396
DI flags	0 (vacuous)	7 (informative)

The flags after mitigation are informative, not a regression — once the model actually discriminates between classes, it surfaces the real disparity the degenerate baseline had hidden (size 24+ DI 0.41; apple/full-bust 0.64–0.70). Height stays equal (≥ 0.99).

Deployed: the class-weighted model. Catching misfits matters more than raw accuracy; the confidence % and blue/amber advisory layer absorb the uncertainty.

Explainability: SHAP on the winning model


TreeExplainer · 3,000 test rows · seed 42



- **size_ordered dominates** (mean |SHAP| 0.38 for small, 0.27 for large vs 0.08 for fit) — and directionally correctly: high ordered size, low weight, small bust band all push toward “runs small”.
- **The model learned real sizing signal;** it is the decision threshold, drowned by class imbalance, that fails — exactly what class weighting attacks.
- **Red flag: missingness is predictive.** weight_kg_missing and body_type_not_reported rank in the top 7 for minority classes. Shoppers who disclose less get systematically different predictions — a data-completeness bias, answered with transparency in the product (confidence drops when key fields are absent).

Architecture: data to deployed product

OFFLINE — reproducible, seed 42



Generative AI: used deliberately, fenced off from the graded ML

Claude API — multimodal fit advice

Measurements + recommendation + the composited try-on image in; two paragraphs of plain-language advice out (the second must start “Note:” with visual observations; tailoring jargon banned in the prompt).

Never estimates measurements · never overrides the model · read-only consumer of the prediction.

rembg / U²-Net — catalog background removal

A pretrained segmentation CNN produces the transparent garment PNGs the try-on compositor warps.

Runs offline, once per image, at catalog build time. Sees no user data, produces no labels, touches no training data.

Why the fence?

- **Auditability** — the audit's quantitative claims need a fixed, inspectable model; a generative model in the prediction path would make them unverifiable
- **No silent data invention** — vision-estimated measurements would inject body-type-correlated error into exactly the variables the audit stratifies on
- **Reproducibility** — graded artifacts rebuild from fixed seeds; generative output is confined to advice prose and one-time asset prep



Limitations and future work

Known limitations (documented, not hidden)

- **Severe class imbalance remains the dominant problem** — one mitigation trialed, fully audited
- Random split does not group repeat customers (IDs dropped upstream)
- Self-reported measurements — noisy, plausibly in group-correlated ways
- Groups are proxies: no age/ethnicity/disability fields in the data
- Men's extension is chart-based and unaudited, by design

Future work

- **Probability calibration + per-group thresholds** to equalize misfit-detection rates — the clear next mitigation
- **FIT dataset (2026, 1.13M fit-aware try-on triplets)** — the research path to photorealistic, fit-accurate generative try-on, upgrading today's size-proportional 2D scaling
- **Transformer clothing parsing** (SegFormer benchmarked: no net win today — 59 vs 56 clean items, 5 failures)
- **Color-suitability matching** — scoped out from the start
- **bust_input_method audit** once real usage data exists



Part III

Business Presentation Deck

Audience: executives — the returns problem, the product, impact and rollout



FitML: A Fairness-Audited Machine Learning System for Sizing and Augmented Reality Try-On in E-Commerce

Pillar 5: Capstone Project

Postgraduate Diploma in Artificial Intelligence and Machine Learning

Debie Galleros · July 2026

A virtual fitting room with size recommendations and tailored advice

Business presentation — live demo: fit-ml.netlify.app



Returns are eating online fashion's margins

35%

of online clothing orders come back — vs about 8% in physical stores

\$100B+

what returns cost retailers every year

#1

cause: sizing uncertainty — shoppers can't try before they buy

Two questions stop a shopper from clicking “buy”:

“Will it fit?”

Every wrong answer is a return — shipped twice, repacked, often discounted.

“How will it look on me?”

If the shopper can't picture it, they don't buy — or they order three sizes to try at home.

FitML answers both — before the purchase

A working prototype, live at fit-ml.netlify.app

1

Set up a profile once

The shopper enters their measurements and uploads one photo. Nothing is guessed from the photo — it is only used to show garments on them.

2


Browse and try on

117 real items with filters, prices, and color options. One tap shows the garment on the shopper's own photo.

3


Get a size and honest advice

A recommended size with a confidence level, plus two short paragraphs of plain-language advice — including what the try-on photo actually shows.



The live site

The virtual fitting room, on the shopper's own photo



Live try-on: recommendation, composite, and advice for a real catalog item

A clear recommendation

“Recommended size: XS — 40% confidence” — the confidence level is shown, never hidden.

The garment on their body

The sweater is placed on the shopper's own photo, scaled and angled to their pose — picking a bigger size renders visibly bigger.

Advice a shopper actually understands

Plain words, no tailor talk: “the length reaches your waist neatly, and the sleeves land well at your wrists.”

Privacy by default

The face is blurred before the photo is even saved — visible right in this screenshot.



The differentiator: sizing that is checked for everyone

Most sizing tools are black boxes. FitML publishes its audit.

Sizing models are usually most accurate for “average” bodies — and quietly worse for everyone else. We measured ours across body types, heights, and size ranges, found exactly that problem, and retrained to fix what matters most.

What we found: the first model told almost every shopper “it fits” — it was most wrong for plus-size shoppers, exactly the people returns hit hardest.

What we did: retrained the model to take misfits seriously and published before/after numbers for every group — the full report is in the open repository.

1 in 100 → 1 in 2

misfit warnings the model catches, before vs after the fairness fix — the single number behind fewer wrong “it fits” promises

8 groups × 3 lenses

body types, heights, and size bands — every group's numbers published, none averaged away

0 guesses

measurements are always entered by the shopper — never estimated from a photo



What this is worth

Projected impact — to be validated in a pilot, not promised in advance

The mechanism

- **Fewer wrong “it fits” moments.** The retrained model catches about half of misfits before purchase, instead of almost none.
- **Fewer “order 3, keep 1” baskets.** Seeing the garment on their own photo replaces ordering multiple sizes to try at home.
- **Higher purchase confidence.** A visible confidence level and honest sizing tips build trust — and trust converts.

Every avoided return saves twice-paid shipping, inspection, repackaging, and markdown loss — and keeps the item sellable at full price.

The target

35% → 15–20%

return-rate goal for sizing-driven returns, in line with what published fit-technology deployments report

Illustrative: a shop doing 10,000 orders a month at 35% returns processes 3,500 returns. Cutting the rate to 20% removes 1,500 of them — every month, before counting the sales that no longer get lost to hesitation.

Risks handled up front

Body photos are sensitive — the prototype treats them that way, by design



Photos deleted in 24 hours

Auto-deletion is enforced in code, not policy. Uploads live in anonymous, session-scoped folders — never linked to a name or email.



Faces blurred by default

The blur happens before the photo is first saved, so an unblurred photo never exists on the server unless the shopper opts in.



No body measurements are ever guessed

The photo is used only to show garments. Measurements come from the shopper, full stop — no silent AI estimation of anyone's body.



Compliant framing from day one

Built against the Philippines' Data Privacy Act (RA 10173): consent up front, one declared purpose, minimal collection, hard retention limits.

And the honest risk register

- The model over-warns today: catching misfits costs some false sizing tips. The interface presents them as tips, not verdicts.
- Trained on US women's rental/retail data; local fit habits will differ. The pilot exists to recalibrate on real local outcomes.
- Men's sizing is a size-chart lookup, clearly labeled — not the audited model.

Rollout: pilot, measure, then scale

Phase 1 — Pilot

8–10 weeks

One women's category (dresses) on the existing store. Measure return rate, size-exchange rate, and conversion against a control group.

Phase 2 — Calibrate

next quarter

Retune the model's thresholds on real local outcomes, per body-type group — the audit framework is already built for exactly this.

Phase 3 — Scale

quarter after

Full women's catalog, then men's once enough real fit feedback exists to train and audit a proper men's model.

The gate between phases is a number, not a feeling: the pilot ships with the same measurement dashboard the capstone already produces — per-group accuracy, misfit-catch rate, and return-rate delta.



Where this goes next

Photorealistic try-on

New research data (the FIT dataset, 2026 — over a million fit-aware try-on examples) makes it possible to render how a specific size would actually drape and pull. Today's preview shows placement and proportion; this upgrade shows fit.

Smarter sizing over time

Every confirmed purchase and return becomes training signal. The fairness dashboard keeps the improvement honest across body types.

Color advice

“Does this shade suit me?” — scoped out of the capstone deliberately, and a natural add-on to the same advice panel.

Richer men's experience

Real men's fit data from the pilot replaces today's clearly-labeled size-chart lookup with an audited model.



Try it yourself

fit-ml.netlify.app

The full fairness audit, code, and documentation are public: github.com/debiegalleros/fitml-capstone

Debie Galleros · Postgraduate Diploma in Artificial Intelligence and Machine Learning · July 2026